

Herding cats with dplyr

Jyotishka Datta

2022-05-23

Contents

dplyr	1
Getting the package	2
Subsetting Data	2
Pipes	7
Combining Select & Filter	8
Going further	11

dplyr

80% of the work involved with data analysis involves cleaning and shaping the data until it's in the state you need. Bracket subsetting is handy, but it can be cumbersome and difficult to read, especially for complicated operations. Enter **dplyr**!

dplyr is a package for making data manipulation easier. (It does a lot more too, but this is what we'll focus on).

Unlike the subsetting commands we've already worked on, **dplyr** is designed to be highly expressive, and highly readable. It's structured around a set of verbs, or grammar of data manipulation. The core functions we'll talk about are below:

- `select`
- `arrange`
- `filter`
- `group_by`
- `mutate`
- `summarise/summarize`

Getting the package

```
install.packages("dplyr")
```

Again, you may be asked to choose a CRAN mirror; RStudio is a good choice.

Unlike the `rmarkdown` package, we'll be using `dplyr` in the console. In order to do that, we need to *load* the package into our environment so we can access functions from `dplyr`. We do this with the `library` command:

```
library("dplyr")
```

You only need to install a package once per computer, but you need to load it every time you open a new R session and want to use that package.

Subsetting Data

The first two `dplyr` commands we'll use help us to subset our data by rows and columns.

`select`

The first command we'll use is `select`, which allows us to choose columns from our dataset. Let's use our `cats` dataset and select only the `coat` column; we did this previously with

```
cats[, "coat"]
```

```
## [1] "tabby"    "maltese" "brown"    "black"    "calico"   "tabby"    "brown"
## [8] "brown"    "black"    "maltese"
```

With `dplyr`, we don't need to enclose our column names in quotes

```
select(cats, coat)
```

```
##      coat
## 1  tabby
## 2 maltese
## 3  brown
## 4  black
## 5  calico
## 6  tabby
## 7  brown
## 8  brown
## 9  black
## 10 maltese
```

Notice how the output differs slightly; all the main `dplyr` verbs behave consistently in that their inputs and outputs are both `data.frames`, rather than returning a simple vector as the bracket-indexing method did. All of the main “verbs” we’ll talk about will return a `data.frame` as their result.

We can select more columns by giving `select` additional arguments, and our output `data.frame` will have columns according to the order of our arguments

```
select(cats, coat, cat_id)
```

```
##      coat cat_id
## 1  tabby   321
## 2 maltese  250
## 3  brown   219
## 4  black   182
## 5  calico  107
## 6  tabby   234
## 7  brown   196
## 8  brown   311
## 9  black   130
## 10 maltese  349
```

`filter`

So where `select` allowed us to select *columns*, `filter` operated on *rows*. Say we want to see the all the cats with black coats; we saw earlier how to use that using bracket-indexing:

```
cats[cats$coat == "black", ]
```

```
##      street  coat  sex  age weight fixed wander_dist roamer cat_id
## 4    140 Robin Way black female 7.172  8.053     1      0.030    no    182
## 9 16982 Kennedy Rd black female 2.851  3.291     0      0.065    no    130
```

In `dplyr`, this looks like

```
filter(cats, coat == "black")
```

```
##      street  coat  sex  age weight fixed wander_dist roamer cat_id
## 1    140 Robin Way black female 7.172  8.053     1      0.030    no    182
## 2 16982 Kennedy Rd black female 2.851  3.291     0      0.065    no    130
```

Notice we don’t have to use the `$` operator to tell `filter` where the `coat` column is; it’s smart enough to assume we want the `coat` column from the `data.frame` we passed in.

arrange

Maybe you have a set of observations in your data that you want to organize by their value. `arrange` allows us to change the order of rows in our dataset based on their values.

```
arrange(cats, coat)
```

```
##           street    coat  sex  age weight fixed wander_dist roamer
## 1      140 Robin Way  black female 7.172  8.053    1      0.030    no
## 2     16982 Kennedy Rd  black female 2.851  3.291    0      0.065    no
## 3      201 Hollywood Ave  brown female 4.601  3.947    1      0.076    no
## 4   115 Via Santa Maria  brown   male 6.917  5.626    1      0.097   yes
## 5      303 Harding Ave  brown   male 3.713  3.982    1      0.033    no
## 6      135 Charles St  calico   male 4.660  6.193    1      0.085   yes
## 7      242 Harding Ave  maltese female 8.234 12.368    1      0.033    no
## 8    16528 Marchmont Dr  maltese female 4.594  6.994    0      0.059    no
## 9       Los Robles Way  tabby  female 3.003  3.993    0      0.040    no
## 10  130 Vista Del Campo  tabby  female 3.796  3.860    1      0.085    no
##   cat_id
## 1     182
## 2     130
## 3     219
## 4     196
## 5     311
## 6     107
## 7     250
## 8     349
## 9     321
## 10    234
```

```
# you can include additional columns to help sort the data
arrange(cats, coat, sex)
```

```
##           street    coat  sex  age weight fixed wander_dist roamer
## 1      140 Robin Way  black female 7.172  8.053    1      0.030    no
## 2     16982 Kennedy Rd  black female 2.851  3.291    0      0.065    no
## 3      201 Hollywood Ave  brown female 4.601  3.947    1      0.076    no
## 4   115 Via Santa Maria  brown   male 6.917  5.626    1      0.097   yes
## 5      303 Harding Ave  brown   male 3.713  3.982    1      0.033    no
## 6      135 Charles St  calico   male 4.660  6.193    1      0.085   yes
## 7      242 Harding Ave  maltese female 8.234 12.368    1      0.033    no
## 8    16528 Marchmont Dr  maltese female 4.594  6.994    0      0.059    no
## 9       Los Robles Way  tabby  female 3.003  3.993    0      0.040    no
## 10  130 Vista Del Campo  tabby  female 3.796  3.860    1      0.085    no
```

```
##      cat_id
## 1      182
## 2      130
## 3      219
## 4      196
## 5      311
## 6      107
## 7      250
## 8      349
## 9      321
## 10     234
```

mutate

One common task in working with data is updating/cleaning some of the values in columns. `mutate` allows us to do this relatively easily. Let's say I don't want a lot of decimal places in one of my measurements. I can use `mutate` to update my existing variable:

```
mutate(cats, weight = round(weight, 2))
```

```
##           street    coat  sex  age weight fixed wander_dist roamer
## 1    Los Robles Way  tabby female 3.003   3.99    0     0.040    no
## 2    242 Harding Ave maltese female 8.234  12.37    1     0.033    no
## 3    201 Hollywood Ave  brown female 4.601   3.95    1     0.076    no
## 4     140 Robin Way   black female 7.172   8.05    1     0.030    no
## 5     135 Charles St  calico   male 4.660   6.19    1     0.085   yes
## 6  130 Vista Del Campo  tabby female 3.796   3.86    1     0.085    no
## 7  115 Via Santa Maria  brown   male 6.917   5.63    1     0.097   yes
## 8     303 Harding Ave  brown   male 3.713   3.98    1     0.033    no
## 9    16982 Kennedy Rd  black female 2.851   3.29    0     0.065    no
## 10  16528 Marchmont Dr maltese female 4.594   6.99    0     0.059    no
##      cat_id
## 1      321
## 2      250
## 3      219
## 4      182
## 5      107
## 6      234
## 7      196
## 8      311
## 9      130
## 10     349
```

Another common task is generating a new column based on values that are already in the dataset you are working on. `mutate` helps us do this, and tacks a new column to the end of our `data.frame`.

let's say you want to add two variables together

```
mutate(cats, new_variable = age + weight)
```

	street	coat	sex	age	weight	fixed	wander_dist	roamer
## 1	Los Robles Way	tabby	female	3.003	3.993	0	0.040	no
## 2	242 Harding Ave	maltese	female	8.234	12.368	1	0.033	no
## 3	201 Hollywood Ave	brown	female	4.601	3.947	1	0.076	no
## 4	140 Robin Way	black	female	7.172	8.053	1	0.030	no
## 5	135 Charles St	calico	male	4.660	6.193	1	0.085	yes
## 6	130 Vista Del Campo	tabby	female	3.796	3.860	1	0.085	no
## 7	115 Via Santa Maria	brown	male	6.917	5.626	1	0.097	yes
## 8	303 Harding Ave	brown	male	3.713	3.982	1	0.033	no
## 9	16982 Kennedy Rd	black	female	2.851	3.291	0	0.065	no
## 10	16528 Marchmont Dr	maltese	female	4.594	6.994	0	0.059	no
##	cat_id	new_variable						
## 1	321	6.996						
## 2	250	20.602						
## 3	219	8.548						
## 4	182	15.225						
## 5	107	10.853						
## 6	234	7.656						
## 7	196	12.543						
## 8	311	7.695						
## 9	130	6.142						
## 10	349	11.588						

you can include as many new variables as you want, separated by a comma

```
mutate(cats, new_var_1 = age + weight, new_var_2 = age * weight)
```

	street	coat	sex	age	weight	fixed	wander_dist	roamer
## 1	Los Robles Way	tabby	female	3.003	3.993	0	0.040	no
## 2	242 Harding Ave	maltese	female	8.234	12.368	1	0.033	no
## 3	201 Hollywood Ave	brown	female	4.601	3.947	1	0.076	no
## 4	140 Robin Way	black	female	7.172	8.053	1	0.030	no
## 5	135 Charles St	calico	male	4.660	6.193	1	0.085	yes
## 6	130 Vista Del Campo	tabby	female	3.796	3.860	1	0.085	no
## 7	115 Via Santa Maria	brown	male	6.917	5.626	1	0.097	yes
## 8	303 Harding Ave	brown	male	3.713	3.982	1	0.033	no
## 9	16982 Kennedy Rd	black	female	2.851	3.291	0	0.065	no
## 10	16528 Marchmont Dr	maltese	female	4.594	6.994	0	0.059	no

```
##      cat_id new_var_1  new_var_2
## 1      321      6.996  11.990979
## 2      250     20.602 101.838112
## 3      219      8.548  18.160147
## 4      182     15.225  57.756116
## 5      107     10.853  28.859380
## 6      234      7.656  14.652560
## 7      196     12.543  38.915042
## 8      311      7.695  14.785166
## 9      130      6.142   9.382641
## 10     349     11.588  32.130436
```

Pipes

You'll often find yourself needing to use multiple functions in a row to organize some data that you're working on. This can sometimes lead to dense code that is difficult to read.

```
# for example
sort(round(sqrt(cats$age * 2), 3))
```

```
## [1] 2.388 2.451 2.725 2.755 3.031 3.033 3.053 3.719 3.787 4.058
```

In the code above, I have multiple steps to get my result, but you have to read what's going on from the inside out. This can be cumbersome, especially if you need to understand how one function's output influences the next operation.

Using Pipes

`dplyr` includes a special operator designed to make code *flow* and appear more readable.

It's written as `%>%`, and you can call it the “pipe” operator.

Our example above can be re-written as:

```
cats$age * 2 %>%
  sqrt() %>%
  round(3) %>%
  sort()
```

```
## [1] 4.246242 11.642876 6.505814 10.141208 6.589240 5.367544 9.780638
## [8] 5.250182 4.031314 6.495916
```

Instead of being nested within a bunch of commands, you can see read the code as a series of statements: 1. Multiply `cats$age` by 2, *then* 2. Take the square-root of these values, *then* 3. Round the result to the 3rd digit, *then* 4. Sort the values in ascending order

I encourage you to think of the `%>%` as short-hand for “then”, when reading code that uses it!

“Pipe” operators are found in other languages; they get their name from the idea that your code can be thought of as a “pipeline”.

Let’s look at another example.

```
round(1.23456789, 3)
```

```
## [1] 1.235
```

We can use a pipe operator to acheive the same thing.

```
1.23456789 %>% round(3)
```

```
## [1] 1.235
```

The pipe takes care of making sure the output of the expression on the left-hand-side (a simple numeric, in this case) is inserted as the first argument of the expressing on the right-hand-side. We can also pipe into other argument positions by using a period as a placeholder.

```
3 %>% round(1.23456789, .)
```

```
## [1] 1.235
```

These are contrived examples, and I don’t suggest using pipes for simple operations like rounding. The pipes really become useful when chaining together multiple operations in sequence, as we’ll do with our `dplyr` functions.

Combining Select & Filter

The pipe is really helpful when combined with the data-manipulation of `dplyr`. Remember how we used `filter` to select only the black cats? What if we only want to see the ID’s of those cats, rather than all the info about them? We’ve already seen we can use `select` to pick out certain columns. We can use that to select the `cat_id` column from our `filtered` dataset like so


```
# reading from the inside out
select(filter(cats, coat == "black"), cat_id)
```

```
##   cat_id
## 1    182
## 2    130
```

That might not look too bad now, but what if we wanted to do another operation on that output? We'd add another layer of nesting, and having to read that line from the inside-out can quickly become annoying. We can use the pipe operator to clean that up.

```
# reading from left to right
filter(cats, coat == "black") %>% select(cat_id)
```

```
##   cat_id
## 1    182
## 2    130
```

We could even add another pipe to feed `cats` into `filter`; it isn't necessary, but it makes it even easier to see what we're operating on in this chain of commands. We'll combine this with some line breaks to really make this easy to read:

```
cats %>%
  filter(coat == "black") %>%
  select(cat_id)
```

```
##   cat_id
## 1    182
## 2    130
```

summarize

While `mutate` creates new columns, it's often useful to summarize multiple rows into a single value. Say we want to find the mean weight of all these cats; enter **summarize**! Like `mutate`, the arguments to `summarize` (after the `data.frame` we want to operate on) are expressions. We can combine `summarize` with the `mean` function to get a mean weight for our collection of cats like so:

```
cats %>% summarize(mean_weight = mean(weight))
```

```
##   mean_weight
## 1         5.8307
```

Notice how we have only a single value returned, but it's still in a `data.frame` format. This is subtle, but important; all these basic `dplyr` verbs take in `data.frames` and also return `data.frames`. This consistency helps make long chains of `dplyr` operations possible.

group_by

A very common data analysis task is to do operations like we did above, but to do them on a group-by-group basis. To do this with `dplyr`, we'll use the `group_by` function.

Let's look at the mean weights of our cats, grouping up by coat. This will give us the mean weight of the black cats, mean weight of the calico cats, etc. We can do this by inserting a `group_by` function into our earlier expression for computing mean weight:

```
cats %>%
  group_by(coat) %>%
  summarize(mean_weight = mean(weight))
```

```
## # A tibble: 5 x 2
##   coat      mean_weight
## * <chr>         <dbl>
## 1 black          5.67
## 2 brown          4.52
## 3 calico         6.19
## 4 maltese        9.68
## 5 tabby          3.93
```

Ta-da!

We can also use `mutate` on a per-group basis. Let's make a new column which centers our weights around zero; this can be done by subtracting the group's mean weight from each cat's weight:

```
cats %>%
  group_by(coat) %>%
  mutate(centered_weight = weight - mean(weight))
```

```
## # A tibble: 10 x 10
## # Groups:   coat [5]
##   street coat sex age weight fixed wander_dist roamer cat_id
##   <chr> <chr> <chr> <dbl> <dbl> <int> <dbl> <chr> <int>
## 1 Los R~ tabby fema~ 3.00 3.99 0 0.04 no 321
## 2 242 H~ malt~ fema~ 8.23 12.4 1 0.033 no 250
## 3 201 H~ brown fema~ 4.60 3.95 1 0.076 no 219
## 4 140 R~ black fema~ 7.17 8.05 1 0.03 no 182
## 5 135 C~ cali~ male 4.66 6.19 1 0.085 yes 107
## 6 130 V~ tabby fema~ 3.80 3.86 1 0.085 no 234
## 7 115 V~ brown male 6.92 5.63 1 0.097 yes 196
## 8 303 H~ brown male 3.71 3.98 1 0.033 no 311
## 9 16982~ black fema~ 2.85 3.29 0 0.065 no 130
## 10 16528~ malt~ fema~ 4.59 6.99 0 0.059 no 349
## # ... with 1 more variable: centered_weight <dbl>
```

Going further

This is an introductory look at `dplyr`, just enough to make you dangerous. As you continue your R journey I suggest looking into the other awesome things you can do with this package!